

# ODL – a Python framework for rapid prototyping in inverse problems

Holger Kohr<sup>1</sup>   Jonas Adler<sup>2</sup>   Ozan Öktem<sup>2</sup>

<sup>1</sup>Computational Imaging  
Centrum Wiskunde & Informatica, Amsterdam

<sup>2</sup>Department of Mathematics  
KTH - Royal Institute of Technology, Stockholm

October 17, 2017  
ODL training  
Max IV, Lund



# Why a new software framework?

- **Multiple modalities:** CT, CBCT, PET, SPECT, spectral CT, phase contrast CT, electron tomography, ...
- **Mathematical structures/notions:** Functional, operator, Fréchet derivative, proximal, diffeomorphism, discretization, sparsifying transforms, ...
- **Flexibility:** Mathematical structures/notions **re-usable across modalities**  
~> Make it easy to “play around” with new ideas and combine concepts.
- **Collaborative research:** Need to share implementations of common concepts
- **Reproducible research:** Not enough to share theory and pseudocode, also need to share data and concrete implementations  
~> Software components need to be usable by others.

# Why a new software framework?

- **Multiple modalities:** CT, CBCT, PET, SPECT, spectral CT, phase contrast CT, electron tomography, ...
- **Mathematical structures/notions:** Functional, operator, Fréchet derivative, proximal, diffeomorphism, discretization, sparsifying transforms, ...
- **Flexibility:** Mathematical structures/notions **re-usable across modalities**  
~> Make it easy to “play around” with new ideas and combine concepts.
- **Collaborative research:** Need to share implementations of common concepts
- **Reproducible research:** Not enough to share theory and pseudocode, also need to share data and concrete implementations  
~> Software components need to be usable by others.

**Conclusion:** Need a **common software framework** to exchange implementations of concepts and methods.

# Why a new software framework?

## Requirements on a software framework:

- Allow formulation and solution of inverse problems in a **common language**.
- Make implementations **re-usable** and **extendable**.
- Enable **fast prototyping** on **clinically relevant data**.
- Leverage the power of **existing libraries**.

# Why a new software framework?

Requirements on a software framework:

- Allow formulation and solution of inverse problems in a **common language**.
- Make implementations **re-usable** and **extendable**.
- Enable **fast prototyping** on **clinically relevant data**.
- Leverage the power of **existing libraries**.

**Initial situation:** No existing framework fit our purpose.

# Operator Discretization Library

An object-oriented Python framework for inverse problems

Main components:

- Functional analysis module

Handling of **vector spaces**, **operators**, **discretizations** – generally with a **continuous** point of view

- Optimization methods module

**General-purpose** optimization methods suitable for solving inverse problems.

- Tomography module

Acquisition **geometries** and **forward operators** for tomographic applications.

# Operator Discretization Library

An object-oriented Python framework for inverse problems

Main components:

- Functional analysis module

Handling of **vector spaces**, **operators**, **discretizations** – generally with a **continuous** point of view

- Optimization methods module

**General-purpose** optimization methods suitable for solving inverse problems.

- Tomography module

Acquisition **geometries** and **forward operators** for tomographic applications.

# Operator Discretization Library

An object-oriented Python framework for inverse problems

Main components:

- Functional analysis module

Handling of **vector spaces**, **operators**, **discretizations** – generally with a **continuous** point of view

- Optimization methods module

**General-purpose** optimization methods suitable for solving inverse problems.

- Tomography module

Acquisition **geometries** and **forward operators** for tomographic applications.



# Operator Discretization Library

An object-oriented Python framework for inverse problems

## Main components:

- Library of atomic mathematical components
  - Deformation operators
  - Function transforms: wavelet, Fourier, ...
  - Differential operators: partial derivative, gradient, Laplacian, ...
  - Discretization-related: (re-)sampling, interpolation, domain extension, ...
- Utility functions
  - Visualization: Slice viewer, real time plotting, ...
  - Phantoms: Shepp-Logan, FORBILD, Defrise, ...
  - Data I/O: MRC2014, Mayo Clinic, ...

# Operator Discretization Library

An object-oriented Python framework for inverse problems

Main components:

- Library of atomic mathematical components
  - Deformation operators
  - Function transforms: wavelet, Fourier, ...
  - Differential operators: partial derivative, gradient, Laplacian, ...
  - Discretization-related: (re-)sampling, interpolation, domain extension, ...
- Utility functions
  - Visualization: Slice viewer, real time plotting, ...
  - Phantoms: Shepp-Logan, FORBILD, Defrise, ...
  - Data I/O: MRC2014, Mayo Clinic, ...

# Operator Discretization Library

An object-oriented Python framework for inverse problems

Main components:

- User-contributed modules

“Fast track” for experimental or slightly exotic code

- Figures of Merit (FOMs) for image quality assessment
- Handlers for specific data formats or geometries
- Functionality to download and import public datasets
- Wrappers for 3 major Deep Learning frameworks: **Tensorflow**, **Theano** and **Pytorch**

# Design principle: modularity

Consider a TV minimization problem

$$\min_{f \in X} \left[ \|\mathcal{T}(f) - g\|_Y^2 + \lambda \text{TV}(f) \right]$$

Components:

- *Reconstruction space  $X$*
- *Data space  $Y$*
- *Forward operator  $\mathcal{T}: X \rightarrow Y$*
- *Data  $g \in Y$*
- *Data discrepancy functional  $\|\cdot - g\|_Y^2$*
- *Regularization parameter  $\lambda > 0$*
- *Regularization functional  $\text{TV}(\cdot)$*

- ↪ (almost) freely exchangeable “modules” in the mathematical formulation
- ↪ ODL maps them to software objects as closely as possible
- ↪ Mathematics as strong guideline for software design
- ↪ Makes the software “feel” natural to mathematicians

# Design principle: modularity

Consider a TV minimization problem

$$\min_{f \in X} \left[ \|\mathcal{T}(f) - g\|_Y^2 + \lambda \text{TV}(f) \right]$$

Components:

- *Reconstruction space  $X$*
- *Data space  $Y$*
- *Forward operator  $\mathcal{T}: X \rightarrow Y$*
- *Data  $g \in Y$*
- *Data discrepancy functional  $\|\cdot - g\|_Y^2$*
- *Regularization parameter  $\lambda > 0$*
- *Regularization functional  $\text{TV}(\cdot)$*

↪ (almost) freely exchangeable “modules” in the mathematical formulation

↪ ODL maps them to software objects as closely as possible

↪ Mathematics as strong guideline for software design

↪ Makes the software “feel” natural to mathematicians

# Design principle: modularity

Consider a TV minimization problem

$$\min_{f \in X} \left[ \|\mathcal{T}(f) - g\|_Y^2 + \lambda \text{TV}(f) \right]$$

Components:

- *Reconstruction space*  $X$
- *Data space*  $Y$
- *Forward operator*  $\mathcal{T}: X \rightarrow Y$
- *Data*  $g \in Y$
- *Data discrepancy functional*  $\|\cdot - g\|_Y^2$
- *Regularization parameter*  $\lambda > 0$
- *Regularization functional*  $\text{TV}(\cdot)$

- ↪ (almost) freely exchangeable “modules” in the mathematical formulation
- ↪ ODL maps them to software objects as closely as possible
- ↪ Mathematics as strong guideline for software design
- ↪ Makes the software “feel” natural to mathematicians

# Design principle: abstraction

**Landweber's method:** Determine  $f$  from given data  $g = \mathcal{T}(f)$  and an initial guess  $f_0$  by

$$f_{k+1} = f_k + \omega \mathcal{T}'(f_k)^*(g - \mathcal{T}(f_k)), \quad k = 0, 1, \dots, K-1$$

Code using ODL:

```
def landweber(T, f, g, omega, K):  
    for i in range(K):  
        f += omega * T.derivative(f).adjoint(g - T(f))
```

- Completely generic (expects operator, data, plus some parameters)
- Uses abstract properties of operators in the iteration:
  - $T(f) \longleftrightarrow \mathcal{T}(f)$  (operator evaluation)
  - $T.derivative(f) \longleftrightarrow \mathcal{T}'(f)$  (derivative *operator* at  $f$ )
  - $T.derivative(f).adjoint \longleftrightarrow \mathcal{T}'(f)^*$  (adjoint of the derivative at  $f$ )

# Design principle: abstraction

Landweber's method: Determine  $f$  from given data  $g = \mathcal{T}(f)$  and an initial guess  $f_0$  by

$$f_{k+1} = f_k + \omega \mathcal{T}'(f_k)^*(g - \mathcal{T}(f_k)), \quad k = 0, 1, \dots, K-1$$

Code using ODL:

```
def landweber(T, f, g, omega, K):  
    for i in range(K):  
        f += omega * T.derivative(f).adjoint(g - T(f))
```

- Completely generic (expects operator, data, plus some parameters)
- Uses abstract properties of operators in the iteration:
  - $T(f) \longleftrightarrow \mathcal{T}(f)$  (operator evaluation)
  - $T.derivative(f) \longleftrightarrow \mathcal{T}'(f)$  (derivative *operator* at  $f$ )
  - $T.derivative(f).adjoint \longleftrightarrow \mathcal{T}'(f)^*$  (adjoint of the derivative at  $f$ )



# Design principle: abstraction

Landweber's method: Determine  $f$  from given data  $g = \mathcal{T}(f)$  and an initial guess  $f_0$  by

$$f_{k+1} = f_k + \omega \mathcal{T}'(f_k)^*(g - \mathcal{T}(f_k)), \quad k = 0, 1, \dots, K-1$$

Code using ODL:

```
def landweber(T, f, g, omega, K):  
    for i in range(K):  
        f += omega * T.derivative(f).adjoint(g - T(f))
```

- Completely generic (expects operator, data, plus some parameters)
- Uses abstract properties of operators in the iteration:
  - $T(f) \longleftrightarrow \mathcal{T}(f)$  (operator evaluation)
  - $T.derivative(f) \longleftrightarrow \mathcal{T}'(f)$  (derivative *operator* at  $f$ )
  - $T.derivative(f).adjoint \longleftrightarrow \mathcal{T}'(f)^*$  (adjoint of the derivative at  $f$ )

# Design principle: abstraction

Landweber's method: Determine  $f$  from given data  $g = \mathcal{T}(f)$  and an initial guess  $f_0$  by

$$f_{k+1} = f_k + \omega \mathcal{T}'(f_k)^*(g - \mathcal{T}(f_k)), \quad k = 0, 1, \dots, K-1$$

Code using ODL:

```
def landweber(T, f, g, omega, K):  
    for i in range(K):  
        f += omega * T.derivative(f).adjoint(g - T(f))
```

- Completely generic (expects operator, data, plus some parameters)
- Uses abstract properties of operators in the iteration:
  - $T(f) \longleftrightarrow \mathcal{T}(f)$  (operator evaluation)
  - $T.derivative(f) \longleftrightarrow \mathcal{T}'(f)$  (derivative *operator* at  $f$ )
  - $T.derivative(f).adjoint \longleftrightarrow \mathcal{T}'(f)^*$  (adjoint of the derivative at  $f$ )

# Design principle: abstraction

Landweber's method: Determine  $f$  from given data  $g = \mathcal{T}(f)$  and an initial guess  $f_0$  by

$$f_{k+1} = f_k + \omega \mathcal{T}'(f_k)^*(g - \mathcal{T}(f_k)), \quad k = 0, 1, \dots, K - 1$$

Code using ODL:

```
def landweber(T, f, g, omega, K):  
    for i in range(K):  
        f += omega * T.derivative(f).adjoint(g - T(f))
```

- $T$  is an Operator that implements a *generic, abstract* interface:  
domain, range, derivative, adjoint, operator *evaluation*
- Lots of tools to build complex operators from simple ones:  
operator arithmetic  $T + S$ , composition  $T * S$ , *product space* operators etc.
- There are *many* readily implement operators in ODL, all implementing the above interface

# Design principle: abstraction

Landweber's method: Determine  $f$  from given data  $g = \mathcal{T}(f)$  and an initial guess  $f_0$  by

$$f_{k+1} = f_k + \omega \mathcal{T}'(f_k)^*(g - \mathcal{T}(f_k)), \quad k = 0, 1, \dots, K - 1$$

Code using ODL:

```
def landweber(T, f, g, omega, K):  
    for i in range(K):  
        f += omega * T.derivative(f).adjoint(g - T(f))
```

- T is an Operator that implements a *generic, abstract* interface: domain, range, derivative, adjoint, operator *evaluation*
- Lots of tools to build complex operators from simple ones: operator arithmetic  $T + S$ , composition  $T * S$ , *product space* operators etc.
- There are *many* readily implement operators in ODL, all implementing the above interface

# Design principle: abstraction

Landweber's method: Determine  $f$  from given data  $g = \mathcal{T}(f)$  and an initial guess  $f_0$  by

$$f_{k+1} = f_k + \omega \mathcal{T}'(f_k)^*(g - \mathcal{T}(f_k)), \quad k = 0, 1, \dots, K - 1$$

Code using ODL:

```
def landweber(T, f, g, omega, K):  
    for i in range(K):  
        f += omega * T.derivative(f).adjoint(g - T(f))
```

- $T$  is an Operator that implements a *generic, abstract* interface: domain, range, derivative, adjoint, operator *evaluation*
- Lots of tools to build complex operators from simple ones: operator arithmetic  $T + S$ , composition  $T * S$ , *product space* operators etc.
- There are *many* readily implement operators in ODL, all implementing the above interface

# Design principle: compartmentalization

- Separates the “what” (abstract interface) of an object class from the “how” (concrete implementation)

**Example:** Fourier transform using NumPy FFT vs. pyFFTW vs. cuFFT

- Allows building generic APIs with the possibility for a new implementation in the future (*extensibility*)

**Example:** L1Norm as a concrete realization of the abstract Functional

- Makes functions and classes individually testable
- Documentation is bundled with the object and immediately visible to the user
- Code becomes more maintainable, low risk for “god objects” or scope glide

# Design principle: compartmentalization

- Separates the “what” (abstract interface) of an object class from the “how” (concrete implementation)  
**Example:** Fourier transform using NumPy FFT vs. pyFFTW vs. cuFFT
- Allows building generic APIs with the possibility for a new implementation in the future (*extensibility*)  
**Example:** L1Norm as a concrete realization of the abstract Functional
- Makes functions and classes individually testable
- Documentation is bundled with the object and immediately visible to the user
- Code becomes more maintainable, low risk for “god objects” or scope glide

# Design principle: compartmentalization

- Separates the “what” (abstract interface) of an object class from the “how” (concrete implementation)  
*Example:* Fourier transform using NumPy FFT vs. pyFFTW vs. cuFFT
- Allows building generic APIs with the possibility for a new implementation in the future (*extensibility*)  
*Example:* L1Norm as a concrete realization of the abstract Functional
- Makes functions and classes individually testable
- Documentation is bundled with the object and immediately visible to the user
- Code becomes more maintainable, low risk for “god objects” or scope glide



# Design principle: compartmentalization

- Separates the “what” (abstract interface) of an object class from the “how” (concrete implementation)  
*Example:* Fourier transform using NumPy FFT vs. pyFFTW vs. cuFFT
- Allows building generic APIs with the possibility for a new implementation in the future (*extensibility*)  
*Example:* L1Norm as a concrete realization of the abstract Functional
- Makes functions and classes individually testable
- Documentation is bundled with the object and immediately visible to the user
- Code becomes more maintainable, low risk for “god objects” or scope glide

# Design principle: compartmentalization

- Separates the “what” (abstract interface) of an object class from the “how” (concrete implementation)  
*Example:* Fourier transform using NumPy FFT vs. pyFFTW vs. cuFFT
- Allows building generic APIs with the possibility for a new implementation in the future (*extensibility*)  
*Example:* L1Norm as a concrete realization of the abstract Functional
- Makes functions and classes individually testable
- Documentation is bundled with the object and immediately visible to the user
- Code becomes more maintainable, low risk for “god objects” or scope glide

# Design compromises

- Distinguish types only if they have *computable* differences.

**Example:** There is only one `FunctionSpace` representing  $\mathcal{F}(\Omega, \mathbb{F}) = \{f: \Omega \rightarrow \mathbb{F}\}$ .

**Reason:** No way to do “interesting” things without symbolic calculus.

- Include only features that *add value* in a specific task.

**Example:** Currently, only nearest neighbor and linear interpolation implemented.

**Reason:** Higher-order schemes are seldom used in imaging, no need so far.

- Enforce mathematical rigor only to help the users. *Don't stand in their way.*

**Example:** For operator composition  $T * S$ , check if `S.range == T.domain`.

Conversely, no convexity check of `Functionals` in optimization methods.

**Reason:** Wrong domains/ranges are a major error source, checking helps users.

Conversely, a convex method may work well even on a non-convex problem.

# Design compromises

- Distinguish types only if they have *computable* differences.

**Example:** There is only one `FunctionSpace` representing  $\mathcal{F}(\Omega, \mathbb{F}) = \{f: \Omega \rightarrow \mathbb{F}\}$ .

**Reason:** No way to do “interesting” things without symbolic calculus.

- Include only features that *add value* in a specific task.

**Example:** Currently, only nearest neighbor and linear interpolation implemented.

**Reason:** Higher-order schemes are seldom used in imaging, no need so far.

- Enforce mathematical rigor only to help the users. *Don't stand in their way.*

**Example:** For operator composition  $T * S$ , check if `S.range == T.domain`.

Conversely, no convexity check of `Functionals` in optimization methods.

**Reason:** Wrong domains/ranges are a major error source, checking helps users.

Conversely, a convex method may work well even on a non-convex problem.

# Design compromises

- Distinguish types only if they have *computable* differences.

**Example:** There is only one `FunctionSpace` representing  $\mathcal{F}(\Omega, \mathbb{F}) = \{f: \Omega \rightarrow \mathbb{F}\}$ .

**Reason:** No way to do “interesting” things without symbolic calculus.

- Include only features that *add value* in a specific task.

**Example:** Currently, only nearest neighbor and linear interpolation implemented.

**Reason:** Higher-order schemes are seldom used in imaging, no need so far.

- Enforce mathematical rigor only to help the users. *Don't stand in their way.*

**Example:** For operator composition  $T * S$ , check if `S.range == T.domain`.

Conversely, no convexity check of `Functionals` in optimization methods.

**Reason:** Wrong domains/ranges are a major error source, checking helps users.

Conversely, a convex method may work well even on a non-convex problem.

## Further design considerations

- ODL is a *prototyping* framework, not a black-box solution
  - ~> Give users freedom to experiment and tinker, do “unorthodox” things
  - ~> Very little “intelligence” that guesses what a user wants
  - ~> Instead: make things “just work” that a typical user would expect to work
- It should be fun to explore the “What if?” scenarios in existing examples
- Make use of external highly optimized code for heavy tasks if adequate
- Don't sacrifice performance!
  - ~> Use libraries in the most efficient way possible (avoid copies, operate in-place, work with low-dimensional arrays, vectorization, broadcasting, ...)
  - ~> Compute on the GPU whenever possible (new fast back-end coming soon)

## Further design considerations

- ODL is a *prototyping* framework, not a black-box solution
  - ↪ Give users freedom to experiment and tinker, do “unorthodox” things
  - ↪ Very little “intelligence” that guesses what a user wants
  - ↪ Instead: make things “just work” that a typical user would expect to work
- It should be fun to explore the “What if?” scenarios in existing examples
- Make use of external highly optimized code for heavy tasks if adequate
- Don't sacrifice performance!
  - ↪ Use libraries in the most efficient way possible (avoid copies, operate in-place, work with low-dimensional arrays, vectorization, broadcasting, ...)
  - ↪ Compute on the GPU whenever possible (new fast back-end coming soon)

## Further design considerations

- ODL is a *prototyping* framework, not a black-box solution
  - ↪ Give users freedom to experiment and tinker, do “unorthodox” things
  - ↪ Very little “intelligence” that guesses what a user wants
  - ↪ Instead: make things “just work” that a typical user would expect to work
- It should be fun to explore the “What if?” scenarios in existing examples
- Make use of external highly optimized code for heavy tasks if adequate
- Don't sacrifice performance!
  - ↪ Use libraries in the most efficient way possible (avoid copies, operate in-place, work with low-dimensional arrays, vectorization, broadcasting, ...)
  - ↪ Compute on the GPU whenever possible (new fast back-end coming soon)



## Further design considerations

- ODL is a *prototyping* framework, not a black-box solution
  - ~> Give users freedom to experiment and tinker, do “unorthodox” things
  - ~> Very little “intelligence” that guesses what a user wants
  - ~> Instead: make things “just work” that a typical user would expect to work
- It should be fun to explore the “What if?” scenarios in existing examples
- Make use of external highly optimized code for heavy tasks if adequate
- Don’t sacrifice performance!
  - ~> Use libraries in the most efficient way possible (avoid copies, operate in-place, work with low-dimensional arrays, vectorization, broadcasting, ...)
  - ~> Compute on the GPU whenever possible (new fast back-end coming soon)

## Example: tomography

**Inverse Problem:** Determine attenuation coefficient  $\mu: \Omega \rightarrow \mathbb{R}$  from its ray transform  $\mathcal{R}: L^2(\Omega) \rightarrow Y$  defined as

$$\mathcal{R}(\mu)(\ell) := \int_{\ell} \mu(x) dx$$

for all lines  $\ell = \{x = t\theta + v \in \mathbb{R}^2 \mid t \in \mathbb{R}, \theta \in S^1, v \perp \theta\}$ .

Given: Noisy data

$$g(\ell) \approx \mathcal{R}(\mu)(\ell)$$

Regularization: Conjugate gradient (CGLS) with early termination

## Example: tomography

**Inverse Problem:** Determine attenuation coefficient  $\mu: \Omega \rightarrow \mathbb{R}$  from its ray transform  $\mathcal{R}: L^2(\Omega) \rightarrow Y$  defined as

$$\mathcal{R}(\mu)(\ell) := \int_{\ell} \mu(x) dx$$

for all lines  $\ell = \{x = t\theta + v \in \mathbb{R}^2 \mid t \in \mathbb{R}, \theta \in S^1, v \perp \theta\}$ .

**Given:** Noisy data

$$g(\ell) \approx \mathcal{R}(\mu)(\ell)$$

**Regularization:** Conjugate gradient (CGLS) with early termination

## Example: tomography

**Inverse Problem:** Determine attenuation coefficient  $\mu: \Omega \rightarrow \mathbb{R}$  from its ray transform  $\mathcal{R}: L^2(\Omega) \rightarrow Y$  defined as

$$\mathcal{R}(\mu)(\ell) := \int_{\ell} \mu(x) dx$$

for all lines  $\ell = \{x = t\theta + v \in \mathbb{R}^2 \mid t \in \mathbb{R}, \theta \in S^1, v \perp \theta\}$ .

**Given:** Noisy data

$$g(\ell) \approx \mathcal{R}(\mu)(\ell)$$

**Regularization:** Conjugate gradient (CGLS) with early termination

# Example: Tomography

## Implementation steps:

- Set up uniformly *discretized* image space  $L^2(\Omega)$  with a rectangular domain  $\Omega$  and  $n_x \times n_y$  pixels
- Create parallel beam geometry with  $P$  angles and  $K$  detector pixels
- Define ray transform  $\mathcal{R}: L^2(\Omega) \rightarrow Y$  (the space  $Y$  is inferred from the geometry)
- Solve inverse problem using CGLS
- Display the results

## Example: Tomography

### Implementation steps:

- Set up uniformly *discretized* image space  $L^2(\Omega)$  with a rectangular domain  $\Omega$  and  $n_x \times n_y$  pixels
- Create parallel beam geometry with  $P$  angles and  $K$  detector pixels
- Define ray transform  $\mathcal{R}: L^2(\Omega) \rightarrow Y$  (the space  $Y$  is inferred from the geometry)
- Solve inverse problem using CGLS
- Display the results

# Example: Tomography

## Implementation steps:

- Set up uniformly *discretized* image space  $L^2(\Omega)$  with a rectangular domain  $\Omega$  and  $n_x \times n_y$  pixels
- Create parallel beam geometry with  $P$  angles and  $K$  detector pixels
- Define ray transform  $\mathcal{R}: L^2(\Omega) \rightarrow Y$  (the space  $Y$  is inferred from the geometry)
- Solve inverse problem using CGLS
- Display the results

# Example: Tomography

## Implementation steps:

- Set up uniformly *discretized* image space  $L^2(\Omega)$  with a rectangular domain  $\Omega$  and  $n_x \times n_y$  pixels
- Create parallel beam geometry with  $P$  angles and  $K$  detector pixels
- Define ray transform  $\mathcal{R}: L^2(\Omega) \rightarrow Y$  (the space  $Y$  is inferred from the geometry)
- Solve inverse problem using CGLS
- Display the results



# Example: Tomography

## Implementation steps:

- Set up uniformly *discretized* image space  $L^2(\Omega)$  with a rectangular domain  $\Omega$  and  $n_x \times n_y$  pixels
- Create parallel beam geometry with  $P$  angles and  $K$  detector pixels
- Define ray transform  $\mathcal{R}: L^2(\Omega) \rightarrow Y$  (the space  $Y$  is inferred from the geometry)
- Solve inverse problem using CGLS
- Display the results

# Example: Tomography

## Code

```
# Create reconstruction space and ray transform
space = odl.uniform_discr([-20, -20], [20, 20], shape=(256, 256))
geometry = odl.tomo.parallel_beam_geometry(space, angles=1000)
ray_transform = odl.tomo.RayTransform(space, geometry)


# Create artificial data with around 5 % noise (data max = 10)
phantom = odl.phantom.shepp_logan(space, modified=True)
g = ray_transform(phantom)
g_noisy = g + 0.5 * odl.phantom.white_noise(ray_transform.range)

# Solve inverse problem
x = space.zero()
odl.solvers.conjugate_gradient_normal(ray_transform, x, g_noisy, niter=20)

# Display results
phantom.show('Phantom')
g_noisy.show('Noisy_data')
x.show('CGLS_after_20_iterations')
```

# Example: Tomography

## Results



## Example: Variable $L^p$ TV denoising

Inverse Problem: Find  $f_\alpha \in L^2(\Omega)$  such that

$$f_\alpha = \arg \min_{f \in L^2(\Omega)} \left[ \|f - g\|_{L^2}^2 + \rho_p(\nabla f) \right]$$

for a given noisy image  $g$ . Here,

$$\rho_p(f) := \int_{\Omega} |f(x)|^{p(x)} dx$$

for given  $p : \Omega \rightarrow [1, 2]$ .

Strategy: Use  $\text{prox}_{\rho_p}$  and a splitting method to find a minimizer [Koh17].



H. K.

Total variation regularization with variable Lebesgue prior.

[arXiv:1702.08807 \[math\]](https://arxiv.org/abs/1702.08807), February 2017.

# Example: Variable $L^p$ TV denoising

Code

```
# Read or generate data, Compute the exponent from the data
...

# Setup functionals and operators
data_matching = odl.solvers.L2NormSquared(reco_space).translated(data)
varlp_func = variable_lp.VariableLpModular.gradient.range, exponent)
regularizer = 2e-1 * varlp_func
constraint = odl.solvers.IndicatorBox(reco_space, -5, 5)


lin_ops = [odl.IdentityOperator(reco_space), gradient]

# Start iteration from the noisy data
x = data.copy()


# Choose optimization parameters and go
odl.solvers.douglas_rachford_pd(x, constraint, [data_matching, regularizer],
                                lin_ops, tau=tau, sigma=sigma, lam=lam,
                                niter=100)
```

# Example: Variable $L^p$ TV denoising


Results



Noisy image



TV denoised



Variable  $L^p$  TV denoised

# Work in progress

- Full multi-indexing support also on the GPU (using Theano's `libgpuarray`)

```
>>> space_cpu = odl.rn((2, 3, 4))  # NumPy backend
>>> x_cpu = space_cpu.zero()
>>> x_cpu[:, :3, ::2] = -5
>>> x_cpu[0]
rn((3, 4)).element(
    [[-5.,  0., -5.,  0.],
     [-5.,  0., -5.,  0.],
     [-5.,  0., -5.,  0.]]
)
```

```
>>> space_gpu = odl.rn((2, 3, 4), impl='gpuarray')  # libgpuarray backend
>>> x_gpu = space_gpu.zero()
>>> x_gpu[:, :3, ::2] = -5
>>> x_gpu[0]
rn((3, 4), impl='gpuarray').element(
    [[-5.,  0., -5.,  0.],
     [-5.,  0., -5.,  0.],
     [-5.,  0., -5.,  0.]]
)
```

# Work in progress

- Vector- and tensor-valued functions, useful for, e.g., multi-channel problems or deformation fields

```
>>> # Vector-valued functions, 2 components
>>> space = odl.uniform_discr(0, 1, 5, dtype=(float, (2,)))
>>> def f(x):
...     return (x + 1, x ** 2 - 1)
>>> f_elem = space.element(f)
>>> f_elem # Discretized using vectorized evaluation
uniform_discr(0.0, 1.0, 5, dtype=('float64', (2,))).element(
    [[ 1.1    1.3    1.5    1.7    1.9 ]
     [-0.99 -0.91 -0.75 -0.51 -0.19]]
)
```



# Work in progress

- Deep integration with Deep Learning frameworks

```
>>> odl_op = odl.RayTransform(...)
>>> # Make a layer for Tensorflow
>>> tf_layer = odl.contrib.tensorflow.as_tensorflow_layer(odl_op)
>>> # Make a Theano operator
>>> theano_op = odl.contrib.theano.TheanoOperator(odl_op)
>>> # Make Torch autograd Function
>>> torch_op = odl.contrib.pytorch.TorchOperator(odl_op)
```

Fully support backpropagation (automatic differentiation) via  
`operator.derivative(x).adjoint`