

ODL

A Python framework for rapid prototyping in inverse problems

Jonas Adler^{1, 2} Holger Kohr³ Ozan Öktem¹

¹Department of Mathematics
KTH - Royal Institute of Technology, Stockholm

²Research and Physics
Elekta, Stockholm

³Thermo Fisher Scientific, Eindhoven



Why a new software framework?

Multiple modalities: CT, CBCT, PET, SPECT, spectral CT, phase contrast CT, electron tomography, MRI, HADAF-STEM ...

Why a new software framework?

Multiple modalities: CT, CBCT, PET, SPECT, spectral CT, phase contrast CT, electron tomography, MRI, HADAF-STEM ...

Collaborative research: Need to share implementations of common concepts

Why a new software framework?

Multiple modalities: CT, CBCT, PET, SPECT, spectral CT, phase contrast CT, electron tomography, MRI, HADAF-STEM ...

Collaborative research: Need to share implementations of common concepts

Reproducible research: Not enough to share theory and pseudocode, also need to share data and concrete implementations
~> Software components need to be usable by others.

Why a new software framework?

Multiple modalities: CT, CBCT, PET, SPECT, spectral CT, phase contrast CT, electron tomography, MRI, HADAF-STEM ...

Collaborative research: Need to share implementations of common concepts

Reproducible research: Not enough to share theory and pseudocode, also need to share data and concrete implementations

~> Software components need to be usable by others.

Flexibility: Mathematical structures/notions *re-usable across modalities*

~> Make it easy to “play around” with new ideas and combine concepts.

Why a new software framework?

Multiple modalities: CT, CBCT, PET, SPECT, spectral CT, phase contrast CT, electron tomography, MRI, HADAF-STEM ...

Collaborative research: Need to share implementations of common concepts

Reproducible research: Not enough to share theory and pseudocode, also need to share data and concrete implementations

~> Software components need to be usable by others.

Flexibility: Mathematical structures/notions *re-usable across modalities*

~> Make it easy to “play around” with new ideas and combine concepts.

Conclusion: Need a *common software framework* to exchange implementations of concepts and methods.

Why a new software framework?

Requirements on a software framework:

Allow formulation and solution of inverse problems in a *common language*.

Why a new software framework?

Requirements on a software framework:

Allow formulation and solution of inverse problems in a *common language*.

Make implementations *re-usable* and *extendable*.

Why a new software framework?

Requirements on a software framework:

Allow formulation and solution of inverse problems in a *common language*.

Make implementations *re-usable* and *extendable*.

Enable *fast prototyping* on *clinically relevant data*.

Why a new software framework?

Requirements on a software framework:

Allow formulation and solution of inverse problems in a *common language*.

Make implementations *re-usable* and *extendable*.

Enable *fast prototyping* on *clinically relevant data*.

Leverage the power of *existing libraries*.

Why a new software framework?

Requirements on a software framework:

Allow formulation and solution of inverse problems in a *common language*.

Make implementations *re-usable* and *extendable*.

Enable *fast prototyping* on *clinically relevant data*.

Leverage the power of *existing libraries*.

Initial situation: No existing framework fit our purpose.

Operator Discretization Library

Main components:

- Functional analysis module

- Handling of *vector spaces*, *operators*, *discretizations* – generally with a *continuous* point of view

Operator Discretization Library

Main components:

- Functional analysis module

- Handling of *vector spaces*, *operators*, *discretizations* – generally with a *continuous* point of view

- Optimization methods module

- General-purpose* optimization methods suitable for solving inverse problems.

Operator Discretization Library

Main components:

- Functional analysis module

- Handling of *vector spaces*, *operators*, *discretizations* – generally with a *continuous* point of view

- Optimization methods module

- General-purpose* optimization methods suitable for solving inverse problems.

- Tomography module

- Acquisition *geometries* and *forward operators* for tomographic applications.

Operator Discretization Library

Main components:

Library of atomic mathematical components

- Deformation operators
- Function transforms: wavelet, Fourier, shearlet, ...
- Differential operators: partial derivative, gradient, Laplacian, ...
- Discretization-related: (re-)sampling, interpolation, domain extension, ...

Operator Discretization Library

Main components:

Library of atomic mathematical components

- Deformation operators
- Function transforms: wavelet, Fourier, shearlet, ...
- Differential operators: partial derivative, gradient, Laplacian, ...
- Discretization-related: (re-)sampling, interpolation, domain extension, ...

Utility functions

- Visualization: Slice viewer, real time plotting, ...
- Phantoms: Shepp-Logan, FORBILD, Defrise, ...
- Data I/O: MRC2014, Mayo Clinic, ...

Operator Discretization Library

Main components:

User-contributed modules

“Fast track” for experimental or slightly exotic code

- Figures of Merit (FOMs) for image quality assessment
- Handlers for specific data formats or geometries
- Functionality to download and import public datasets
- Wrappers for Deep Learning frameworks: Tensorflow, Theano, Pytorch, ...

Design principle: modularity

Consider a TV minimization problem

$$\min_{f \in X} [\|\mathcal{T}(f) - g\|_Y^2 + \lambda \text{TV}(f)]$$

Components:

Reconstruction space X

Design principle: modularity

Consider a TV minimization problem

$$\min_{f \in X} [\|\mathcal{T}(f) - g\|_Y^2 + \lambda \text{TV}(f)]$$

Components:

Reconstruction space X

Data space Y

Design principle: modularity

Consider a TV minimization problem

$$\min_{f \in X} [\|\mathcal{T}(f) - g\|_Y^2 + \lambda \text{TV}(f)]$$

Components:

Reconstruction space X

Data space Y

Forward operator $\mathcal{T}: X \rightarrow Y$

Design principle: modularity

Consider a TV minimization problem

$$\min_{f \in X} [\|\mathcal{T}(f) - g\|_Y^2 + \lambda \text{TV}(f)]$$

Components:

Reconstruction space X

Data space Y

Forward operator $\mathcal{T}: X \rightarrow Y$

Data $g \in Y$

Design principle: modularity

Consider a TV minimization problem

$$\min_{f \in X} [\|\mathcal{T}(f) - g\|_Y^2 + \lambda \text{TV}(f)]$$

Components:

Reconstruction space X

Data discrepancy functional $\|\cdot - g\|_Y^2$

Data space Y

Forward operator $\mathcal{T}: X \rightarrow Y$

Data $g \in Y$

Design principle: modularity

Consider a TV minimization problem

$$\min_{f \in X} [\|\mathcal{T}(f) - g\|_Y^2 + \lambda \text{TV}(f)]$$

Components:

Reconstruction space X

Data space Y

Forward operator $\mathcal{T}: X \rightarrow Y$

Data $g \in Y$

Data discrepancy functional $\|\cdot - g\|_Y^2$

Regularization parameter $\lambda > 0$

Design principle: modularity

Consider a TV minimization problem

$$\min_{f \in X} [\|\mathcal{T}(f) - g\|_Y^2 + \lambda \text{TV}(f)]$$

Components:

Reconstruction space X

Data space Y

Forward operator $\mathcal{T}: X \rightarrow Y$

Data $g \in Y$

Data discrepancy functional $\|\cdot - g\|_Y^2$

Regularization parameter $\lambda > 0$

Regularization functional $\text{TV}(\cdot)$

Design principle: modularity

Consider a TV minimization problem

$$\min_{f \in X} [\|\mathcal{T}(f) - g\|_Y^2 + \lambda \text{TV}(f)]$$

Components:

Reconstruction space X

Data space Y

Forward operator $\mathcal{T}: X \rightarrow Y$

Data $g \in Y$

Data discrepancy functional $\|\cdot - g\|_Y^2$

Regularization parameter $\lambda > 0$

Regularization functional $\text{TV}(\cdot)$

\rightsquigarrow (almost) freely exchangeable “modules” in the mathematical formulation

Design principle: modularity

Consider a TV minimization problem

$$\min_{f \in X} [\|\mathcal{T}(f) - g\|_Y^2 + \lambda \text{TV}(f)]$$

Components:

Reconstruction space X

Data space Y

Forward operator $\mathcal{T}: X \rightarrow Y$

Data $g \in Y$

Data discrepancy functional $\|\cdot - g\|_Y^2$

Regularization parameter $\lambda > 0$

Regularization functional $\text{TV}(\cdot)$

\rightsquigarrow (almost) freely exchangeable “modules” in the mathematical formulation

\rightsquigarrow ODL maps them to software objects as closely as possible

Design principle: modularity

Consider a TV minimization problem

$$\min_{f \in X} [\|\mathcal{T}(f) - g\|_Y^2 + \lambda \text{TV}(f)]$$

Components:

Reconstruction space X

Data space Y

Forward operator $\mathcal{T}: X \rightarrow Y$

Data $g \in Y$

Data discrepancy functional $\|\cdot - g\|_Y^2$

Regularization parameter $\lambda > 0$

Regularization functional $\text{TV}(\cdot)$

↪ (almost) freely exchangeable “modules” in the mathematical formulation

↪ ODL maps them to software objects as closely as possible

↪ Mathematics as strong guideline for software design

Design principle: modularity

Consider a TV minimization problem

$$\min_{f \in X} [\|\mathcal{T}(f) - g\|_Y^2 + \lambda \text{TV}(f)]$$

Components:

Reconstruction space X

Data space Y

Forward operator $\mathcal{T}: X \rightarrow Y$

Data $g \in Y$

Data discrepancy functional $\|\cdot - g\|_Y^2$

Regularization parameter $\lambda > 0$

Regularization functional $\text{TV}(\cdot)$

- ~> (almost) freely exchangeable “modules” in the mathematical formulation
- ~> ODL maps them to software objects as closely as possible
- ~> Mathematics as strong guideline for software design
- ~> **Makes the software “feel” natural**

Design principle: abstraction

Landweber's method: Determine f from given data $g = \mathcal{T}(f)$ and initial guess f_0 by

$$f_{k+1} = f_k + \omega [\partial \mathcal{T}(f_k)]^* (g - \mathcal{T}(f_k)), \quad k = 0, 1, \dots, K-1$$

Design principle: abstraction

Landweber's method: Determine f from given data $g = \mathcal{T}(f)$ and initial guess f_0 by

$$f_{k+1} = f_k + \omega [\partial \mathcal{T}(f_k)]^* (g - \mathcal{T}(f_k)), \quad k = 0, 1, \dots, K-1$$

```
def landweber(T, f, g, omega, K):  
    for i in range(K):  
        f += omega * T.derivative(f).adjoint(g - T(f))
```

Design principle: abstraction

Landweber's method: Determine f from given data $g = \mathcal{T}(f)$ and initial guess f_0 by

$$f_{k+1} = f_k + \omega [\partial \mathcal{T}(f_k)]^* (g - \mathcal{T}(f_k)), \quad k = 0, 1, \dots, K-1$$

```
def landweber(T, f, g, omega, K):  
    for i in range(K):  
        f += omega * T.derivative(f).adjoint(g - T(f))
```

Completely generic (expects operator, data, plus some parameters)

Design principle: abstraction

Landweber's method: Determine f from given data $g = \mathcal{T}(f)$ and initial guess f_0 by

$$f_{k+1} = f_k + \omega [\partial \mathcal{T}(f_k)]^* (g - \mathcal{T}(f_k)), \quad k = 0, 1, \dots, K-1$$

```
def landweber(T, f, g, omega, K):  
    for i in range(K):  
        f += omega * T.derivative(f).adjoint(g - T(f))
```

Completely generic (expects operator, data, plus some parameters)

Uses abstract properties of operators in the iteration:

$\rightsquigarrow T(f) \longleftrightarrow \mathcal{T}(f)$ (operator evaluation)

$\rightsquigarrow T.derivative(f) \longleftrightarrow \partial \mathcal{T}(f)$ (derivative *operator* at f)

$\rightsquigarrow T.derivative(f).adjoint \longleftrightarrow [\partial \mathcal{T}(f)]^*$ (adjoint of the derivative at f)

Design principle: abstraction

Landweber's method: Determine f from given data $g = \mathcal{T}(f)$ and initial guess f_0 by

$$f_{k+1} = f_k + \omega [\partial \mathcal{T}(f_k)]^* (g - \mathcal{T}(f_k)), \quad k = 0, 1, \dots, K-1$$

```
def landweber(T, f, g, omega, K):  
    for i in range(K):  
        f += omega * T.derivative(f).adjoint(g - T(f))
```

T is an Operator that implements a *generic, abstract* interface:
domain, range, derivative, adjoint, operator *evaluation*

Design principle: abstraction

Landweber's method: Determine f from given data $g = \mathcal{T}(f)$ and initial guess f_0 by

$$f_{k+1} = f_k + \omega [\partial \mathcal{T}(f_k)]^* (g - \mathcal{T}(f_k)), \quad k = 0, 1, \dots, K-1$$

```
def landweber(T, f, g, omega, K):  
    for i in range(K):  
        f += omega * T.derivative(f).adjoint(g - T(f))
```

T is an Operator that implements a *generic, abstract* interface:
domain, range, derivative, adjoint, operator *evaluation*

Lots of tools to build complex operators from simple ones:
operator arithmetic $T + S$, composition $T * S$, product space operators etc.

Design principle: abstraction

Landweber's method: Determine f from given data $g = \mathcal{T}(f)$ and initial guess f_0 by

$$f_{k+1} = f_k + \omega [\partial \mathcal{T}(f_k)]^* (g - \mathcal{T}(f_k)), \quad k = 0, 1, \dots, K-1$$

```
def landweber(T, f, g, omega, K):  
    for i in range(K):  
        f += omega * T.derivative(f).adjoint(g - T(f))
```

T is an Operator that implements a *generic, abstract* interface:
domain, range, derivative, adjoint, operator *evaluation*

Lots of tools to build complex operators from simple ones:
operator arithmetic $T + S$, composition $T * S$, product space operators etc.

There are *many* readily implement operators in ODL, all implementing the above interface

Design principle: compartmentalization

Separates the “what” (abstract interface) of an object class from the “how” (concrete implementation)

Example: Fourier transform using NumPy FFT vs. pyFFTW vs. cuFFT

Design principle: compartmentalization

Separates the “what” (abstract interface) of an object class from the “how” (concrete implementation)

Example: Fourier transform using NumPy FFT vs. pyFFTW vs. cuFFT

Allows building generic APIs with the possibility for a new implementation in the future (*extensibility*)

Example: L1Norm as a concrete realization of the abstract Functional

Design principle: compartmentalization

Separates the “what” (abstract interface) of an object class from the “how” (concrete implementation)

Example: Fourier transform using NumPy FFT vs. pyFFTW vs. cuFFT

Allows building generic APIs with the possibility for a new implementation in the future (*extensibility*)

Example: L1Norm as a concrete realization of the abstract Functional

Makes functions and classes individually testable

Design principle: compartmentalization

Separates the “what” (abstract interface) of an object class from the “how” (concrete implementation)

Example: Fourier transform using NumPy FFT vs. pyFFTW vs. cuFFT

Allows building generic APIs with the possibility for a new implementation in the future (*extensibility*)

Example: L1Norm as a concrete realization of the abstract `Functional`

Makes functions and classes individually testable

Documentation is bundled with the object and immediately visible to the user

Design principle: compartmentalization

Separates the “what” (abstract interface) of an object class from the “how” (concrete implementation)

Example: Fourier transform using NumPy FFT vs. pyFFTW vs. cuFFT

Allows building generic APIs with the possibility for a new implementation in the future (*extensibility*)

Example: L1Norm as a concrete realization of the abstract Functional

Makes functions and classes individually testable

Documentation is bundled with the object and immediately visible to the user

Code becomes more maintainable, low risk for “god objects” or scope glide

Further design considerations

ODL is a *prototyping* framework, not a black-box solution

- ~> Give users freedom to experiment and tinker, do “unorthodox” things
- ~> Very little “intelligence” that guesses what a user wants
- ~> Instead: make things “just work” that a typical user would expect to work

Further design considerations

ODL is a *prototyping* framework, not a black-box solution

~> Give users freedom to experiment and tinker, do “unorthodox” things

~> Very little “intelligence” that guesses what a user wants

~> Instead: make things “just work” that a typical user would expect to work

It should be *fun* to explore the “What if?” scenarios in existing examples

Further design considerations

ODL is a *prototyping* framework, not a black-box solution

~> Give users freedom to experiment and tinker, do “unorthodox” things

~> Very little “intelligence” that guesses what a user wants

~> Instead: make things “just work” that a typical user would expect to work

It should be *fun* to explore the “What if?” scenarios in existing examples

Make use of external highly optimized code for heavy tasks if adequate

Further design considerations

ODL is a *prototyping* framework, not a black-box solution

~> Give users freedom to experiment and tinker, do “unorthodox” things

~> Very little “intelligence” that guesses what a user wants

~> Instead: make things “just work” that a typical user would expect to work

It should be *fun* to explore the “What if?” scenarios in existing examples

Make use of external highly optimized code for heavy tasks if adequate

Don't sacrifice performance!

~> Use libraries in the most efficient way possible (avoid copies, operate in-place, work with low-dimensional arrays, vectorization, broadcasting, ...)

~> Compute on the GPU whenever possible (new fast back-end coming soon)

Example: Tomography

Inverse Problem: Determine attenuation coefficient $\mu: \Omega \rightarrow \mathbb{R}$ from its ray transform $\mathcal{P}: L^2(\Omega) \rightarrow Y$ defined as

$$\mathcal{P}(\mu)(\ell) := \int_{\ell} \mu(\mathbf{x}) d\mathbf{x}$$

for all lines ℓ .

Example: Tomography

Inverse Problem: Determine attenuation coefficient $\mu: \Omega \rightarrow \mathbb{R}$ from its ray transform $\mathcal{P}: L^2(\Omega) \rightarrow Y$ defined as

$$\mathcal{P}(\mu)(\ell) := \int_{\ell} \mu(\mathbf{x}) d\mathbf{x}$$

for all lines ℓ .

Given: Noisy data

$$g(\ell) \approx \mathcal{P}(\mu)(\ell)$$

Example: Tomography

Inverse Problem: Determine attenuation coefficient $\mu: \Omega \rightarrow \mathbb{R}$ from its ray transform $\mathcal{P}: L^2(\Omega) \rightarrow Y$ defined as

$$\mathcal{P}(\mu)(\ell) := \int_{\ell} \mu(\mathbf{x}) d\mathbf{x}$$

for all lines ℓ .

Given: Noisy data

$$g(\ell) \approx \mathcal{P}(\mu)(\ell)$$

Regularization: Conjugate gradient (CGLS) with early termination

Example: Tomography

Implementation steps:

Set up uniformly *discretized* image space $L^2(\Omega)$ with a rectangular domain Ω and $n_x \times n_y$ pixels

Example: Tomography

Implementation steps:

Set up uniformly *discretized* image space $L^2(\Omega)$ with a rectangular domain Ω and $n_x \times n_y$ pixels

Create parallel beam geometry with P angles and K detector pixels

Example: Tomography

Implementation steps:

Set up uniformly *discretized* image space $L^2(\Omega)$ with a rectangular domain Ω and $n_x \times n_y$ pixels

Create parallel beam geometry with P angles and K detector pixels

Define ray transform $\mathcal{P}: L^2(\Omega) \rightarrow Y$ (the space Y is inferred from the geometry)

Example: Tomography

Implementation steps:

Set up uniformly *discretized* image space $L^2(\Omega)$ with a rectangular domain Ω and $n_x \times n_y$ pixels

Create parallel beam geometry with P angles and K detector pixels

Define ray transform $\mathcal{P}: L^2(\Omega) \rightarrow Y$ (the space Y is inferred from the geometry)

Solve inverse problem using CGLS

Example: Tomography

Implementation steps:

Set up uniformly *discretized* image space $L^2(\Omega)$ with a rectangular domain Ω and $n_x \times n_y$ pixels

Create parallel beam geometry with P angles and K detector pixels

Define ray transform $\mathcal{P}: L^2(\Omega) \rightarrow Y$ (the space Y is inferred from the geometry)

Solve inverse problem using CGLS

Display the results

Example: Tomography




```
# Create reconstruction space and ray transform
space = odl.uniform_discr([-20, -20], [20, 20], shape=(256, 256))
geometry = odl.tomo.parallel_beam_geometry(space, num_angles=1000)
ray_transform = odl.tomo.RayTransform(space, geometry)

# Create artificial data with around 5 % noise (data max = 10)
phantom = odl.phantom.shepp_logan(space, modified=True)
g = ray_transform(phantom)
g_noisy = g + 0.5 * odl.phantom.white_noise(ray_transform.range)

# Solve inverse problem
x = space.zero()
odl.solvers.conjugate_gradient_normal(ray_transform, x, g_noisy, niter=20)

# Display results
phantom.show('Phantom')
g_noisy.show('Noisy data')
x.show('CGLS after 20 iterations')
```

Example: Tomography



Conclusions and Outlook

Reproducible – scalable research requires rethinking scientific software

`github.com/odlgroup/odl`

Conclusions and Outlook

Reproducible – scalable research requires rethinking scientific software
ODL aims to solve these problems

`github.com/odlgroup/odl`

Conclusions and Outlook

Reproducible – scalable research requires rethinking scientific software

ODL aims to solve these problems

Well tested, used in ≈ 20 articles the last year

`github.com/odlgroup/odl`

Conclusions and Outlook

Reproducible – scalable research requires rethinking scientific software

ODL aims to solve these problems

Well tested, used in ≈ 20 articles the last year

Anyone is welcome to use and/or contribute!

`github.com/odlgroup/odl`