

# Learning to Reconstruct Medical Images

Jonas Adler<sup>1, 2</sup>   Ozan Öktem<sup>1</sup>

<sup>1</sup>Department of Mathematics  
KTH - Royal Institute of Technology, Stockholm

<sup>2</sup>Research and Physics  
Elekta, Stockholm



# Why do we need machine learning?

Task: Identify a rabbit in an image

# Why do we need machine learning?

Task: Identify a rabbit in an image

Proposed solution: If the animal is within this range of colors and has long ears and fur and has a slightly elliptical shape and has a nose like... then it is a rabbit

# Why do we need machine learning?

Task: Identify a rabbit in an image


Proposed solution: If the animal is within this range of colors and has long ears and fur and has a slightly elliptical shape and has a nose like... then it is a rabbit



# Inverse Problems

$$g = \mathcal{T}(f_{\text{true}}) + \delta g.$$



- ▶  $g \in Y$       *Data*
- ▶  $f_{\text{true}} \in X$       Image
- ▶  $\mathcal{T} : X \rightarrow Y$       Forward operator
- ▶  $\delta g \in Y$       Noise



# Inverse Problems

$$g = \mathcal{T}(f_{\text{true}}) + \delta g.$$


- ▶  $g \in Y$  Data
- ▶  $f_{\text{true}} \in X$  Image
- ▶  $\mathcal{T} : X \rightarrow Y$  Forward operator
- ▶  $\delta g \in Y$  Noise




# Inverse Problems

$$g = \mathcal{T}(f_{\text{true}}) + \delta g.$$

- ▶  $g \in Y$  Data
- ▶  $f_{\text{true}} \in X$  Image
- ▶  $\mathcal{T} : X \rightarrow Y$  *Forward operator*
- ▶  $\delta g \in Y$  Noise




$\mathcal{T}$   
→




# Inverse Problems

$$g = \mathcal{T}(f_{\text{true}}) + \delta g.$$

- ▶  $g \in Y$  Data
- ▶  $f_{\text{true}} \in X$  Image
- ▶  $\mathcal{T} : X \rightarrow Y$  Forward operator
- ▶  $\delta g \in Y$  *Noise*



$\mathcal{T}$   
 $\rightarrow$







# Inverse Problems

$$g = \mathcal{T}(f_{\text{true}}) + \delta g.$$

- ▶  $g \in Y$  Data
- ▶  $f_{\text{true}} \in X$  Image
- ▶  $\mathcal{T} : X \rightarrow Y$  Forward operator
- ▶  $\delta g \in Y$  Noise



$\mathcal{T}$   
 $\rightarrow$   
 $\leftarrow$   
"  $\mathcal{T}^{-1}$  "



# Solution methods

- ▶ Analytic pseudoinverse (FBP, FDK)  
 $f = \mathcal{T}^\dagger(g)$

# Solution methods

- ▶ Analytic pseudoinverse (FBP, FDK)  
 $f = \mathcal{T}^\dagger(g)$
- ▶ *Variational methods (TV, TGV, Huber)*  
 $f = \arg \min_f ||\mathcal{T}(f) - g||_Y^2 + \lambda ||\nabla f||_1$

# Variational methods

- Strategy, solve an optimization problem:

$$f = \arg \min_f ||\mathcal{T}(f) - \mathbf{g}||_Y^2 + \lambda ||\nabla f||_1$$

Several issues:

# Variational methods

- Strategy, solve an optimization problem:

$$f = \arg \min_f ||\mathcal{T}(f) - g||_Y^2 + \lambda ||\nabla f||_1$$

Several issues:

- Prior is typically unknown - have to "guess"

# Variational methods

- Strategy, solve an optimization problem:

$$f = \arg \min_f ||\mathcal{T}(f) - g||_Y^2 + \lambda ||\nabla f||_1$$

Several issues:

- Prior is typically unknown - have to "guess"
- Parameters ( $\lambda$ ) need to be selected

# Variational methods

- Strategy, solve an optimization problem:

$$f = \arg \min_f ||\mathcal{T}(f) - g||_Y^2 + \lambda ||\nabla f||_1$$

Several issues:

- Prior is typically unknown - have to "guess"
- Parameters ( $\lambda$ ) need to be selected
- Large computational burden

# Solution methods

- ▶ Analytic pseudoinverse (FBP, FDK)

$$f = \mathcal{T}^\dagger(g)$$

- ▶ Variational methods (TV, TGV, Huber)

$$f = \arg \min_f ||\mathcal{T}(f) - g||_\gamma^2 + \lambda ||\nabla f||_1$$

- ▶ *Machine learning*

$$f = \mathcal{T}_\theta^\dagger(g)$$



# Supervised learning

- ▶ We are given training data  $(f, g)$  which is a  $X \times Y$  valued random variable such that  $\mathcal{T}(f) \approx g$ .

# Supervised learning

- ▶ We are given training data  $(f, g)$  which is a  $X \times Y$  valued random variable such that  $\mathcal{T}(f) \approx g$ .
- ▶ We give a class of operators  $\mathcal{T}_\theta^\dagger: Y \rightarrow X$

# Supervised learning

- ▶ We are given training data  $(f, g)$  which is a  $X \times Y$  valued random variable such that  $\mathcal{T}(f) \approx g$ .
- ▶ We give a class of operators  $\mathcal{T}_\theta^\dagger: Y \rightarrow X$
- ▶ Parametrized by  $\theta$  which we *learn*

# Supervised learning

- ▶ We are given training data  $(f, g)$  which is a  $X \times Y$  valued random variable such that  $\mathcal{T}(f) \approx g$ .
- ▶ We give a class of operators  $\mathcal{T}_\theta^\dagger: Y \rightarrow X$
- ▶ Parametrized by  $\theta$  which we learn
- ▶ Selected by optimization of a *loss* function  $L(\theta)$

$$\theta^* = \arg \min_{\theta} L(\theta)$$

# Supervised learning

- ▶ We are given training data  $(f, g)$  which is a  $X \times Y$  valued random variable such that  $\mathcal{T}(f) \approx g$ .
- ▶ We give a class of operators  $\mathcal{T}_\theta^\dagger: Y \rightarrow X$
- ▶ Parametrized by  $\theta$  which we learn
- ▶ Selected by optimization of a *loss* function  $L(\theta)$

$$\theta^* = \arg \min_{\theta} L(\theta)$$


- ▶ *Different from classification ( $X \rightarrow \mathbb{R}^n$ ) and image processing ( $X \rightarrow X$ )*

# Learned inversion methods


- ▶ *Fully learned*
- ▶ Learned post-processing
- ▶ Learned iterative schemes

# Fully learned reconstruction


Goal: Learn "the whole" mapping from data to signal

 *Tomographic image reconstruction based on artificial neural network (ANN) techniques*

Argyrou et. al. NSS/MIC 2012

 *Tomographic image reconstruction using artificial neural networks.*


Paschalis et. al. Nucl Instrum Methods Phys Res A 2004

 *Image reconstruction by domain-transform manifold learning.*


Zhu et. al. Nature 2018

# Fully learned reconstruction


Goal: Learn "the whole" mapping from data to signal

 *Tomographic image reconstruction based on artificial neural network (ANN) techniques*

Argyrou et. al. NSS/MIC 2012

 *Tomographic image reconstruction using artificial neural networks.*

Paschalis et. al. Nucl Instrum Methods Phys Res A 2004

 *Image reconstruction by domain-transform manifold learning.*

Zhu et. al. Nature 2018

Problem:  $\mathcal{T}$  typically has symmetries, but the network has to learn them.

Example: 3D CBCT, data:  $10^8$  pixels and  $10^8$  voxels  $\implies 10^{16}$  connections!



# Learned inversion methods

- ▶ Fully learned
- ▶ *Learned post-processing*
- ▶ Learned iterative schemes

# Learned post-processing

Use deep learning to improve the result of another reconstruction

$$\mathcal{T}_\theta^\dagger = \Lambda_\theta \circ \mathcal{T}^\dagger$$

where  $\mathcal{T}^\dagger$  is some reconstruction (FBP, TV, ...) and  $\Lambda_\theta$  is a learned post-processing operator.

# Learned post-processing

Use deep learning to improve the result of another reconstruction

$$\mathcal{T}_\theta^\dagger = \Lambda_\theta \circ \mathcal{T}^\dagger$$

where  $\mathcal{T}^\dagger$  is some reconstruction (FBP, TV, ...) and  $\Lambda_\theta$  is a learned post-processing operator.


Allows *separation of inversion and learning*, data can be seen as  $(\underbrace{\mathcal{T}^\dagger(\mathbf{g})}_{\in X}, \underbrace{\mathbf{f}}_{\in X})$ .

The problem becomes an image processing problem  $\implies$  easy to solve.

# Learned post-processing

Denoise in transform domain (Fourier, Wavelet, Shearlet, etc)

Won AAPM Low-Dose CT Grand Challenge:

 *A deep convolutional neural network using directional wavelets for low-dose X-ray CT reconstruction*  
Kang et. al. 2016

# Information theoretic bottleneck

►  $I(g) \leftarrow$  information in  $g$

# Information theoretic bottleneck

- ▶  $I(g) \leftarrow$  information in  $g$
- ▶  $I_{total} = I_{prior} + I(g)$

# Information theoretic bottleneck

- ▶  $I(g) \leftarrow$  information in  $g$
- ▶  $I_{total} = I_{prior} + I(g)$
- ▶ For post processing, we only have  $I(\mathcal{T}^\dagger(g))$

# Information theoretic bottleneck

- ▶  $I(g) \leftarrow$  information in  $g$
- ▶  $I_{total} = I_{prior} + I(g)$
- ▶ For post processing, we only have  $I(\mathcal{T}^\dagger(g))$
- ▶  $I(\mathcal{T}^\dagger(g)) \leq I(g)$



# Information theoretic bottleneck

- ▶  $I(g) \leftarrow$  information in  $g$
- ▶  $I_{total} = I_{prior} + I(g)$
- ▶ For post processing, we only have  $I(\mathcal{T}^\dagger(g))$
- ▶  $I(\mathcal{T}^\dagger(g)) \leq I(g)$
- ▶ Post processing beats traditional methods because it utilizes the prior information better

# Information theoretic bottleneck

- ▶  $I(g) \leftarrow$  information in  $g$
- ▶  $I_{total} = I_{prior} + I(g)$
- ▶ For post processing, we only have  $I(\mathcal{T}^\dagger(g))$
- ▶  $I(\mathcal{T}^\dagger(g)) \leq I(g)$
- ▶ Post processing beats traditional methods because it utilizes the prior information better

We could do better by using the raw data!

# Learned inversion methods

- ▶ Fully learned
- ▶ Learned post-processing
- ▶ *Learned iterative schemes*

# Learned iterative reconstruction

- Problem: Data  $g \in Y$ , reconstruction  $f \in X$   
How to include data in each iteration?

# Learned iterative reconstruction

- ▶ Problem: Data  $g \in Y$ , reconstruction  $f \in X$   
How to include data in each iteration?
- ▶ Inspiration from iterative optimization methods

$$f = \arg \min \frac{1}{2} \| \mathcal{T}(f) - g \|_Y^2$$

---

## Algorithm 1 Generic gradient based optimization algorithm

---

```
1: for  $i = 1, \dots$  do  
2:    $f_{i+1} \leftarrow \text{Update}(f_i, \mathcal{T}^*(\mathcal{T}(f_i) - g))$ 
```

---

Gradient descent:

$$\text{Update}(f_i, \mathcal{T}^*(\mathcal{T}(f_i) - g)) = f_i - \alpha \mathcal{T}^*(\mathcal{T}(f_i) - g)$$

# Learned gradient descent

- ▶ Set a stopping criteria (fixed number of steps)
- ▶ Learn the function  $\text{Update} = \Lambda_\theta$

# Learned gradient descent

- ▶ Set a stopping criteria (fixed number of steps)
- ▶ Learn the function  $\text{Update} = \Lambda_\theta$

---

**Algorithm 2** Learned gradient descent

---

```
1: for  $i = 1, \dots, l$  do  
2:    $f_{i+1} \leftarrow \Lambda_\theta(f_i, \mathcal{T}^*(\mathcal{T}(f_i) - g))$   
3:  $\mathcal{T}_\theta^\dagger(g) \leftarrow f_l$ 
```

---

# Learned gradient descent

- ▶ Set a stopping criteria (fixed number of steps)
- ▶ Learn the function  $\text{Update} = \Lambda_\theta$

---

**Algorithm 2** Learned gradient descent

---

```
1: for  $i = 1, \dots, l$  do  
2:    $f_{i+1} \leftarrow \Lambda_\theta(f_i, \mathcal{T}^*(\mathcal{T}(f_i) - g))$   
3:  $\mathcal{T}_\theta^\dagger(g) \leftarrow f_l$ 
```

---

We separate problem dependent (and possibly global) components into  $\mathcal{T}^*(\mathcal{T}(f_i) - g)$ , and local into  $\Lambda_\theta$ !



# Learned Primal-Dual

---

**Algorithm 3** Learned Primal-Dual (conceptual)

---

```
1: for  $i = 1, \dots, l$  do  
2:    $h_i \leftarrow \Gamma_{\theta_i^d}(h_{i-1}, \mathcal{T}(f_{i-1}), g)$   
3:    $f_i \leftarrow \Lambda_{\theta_i^p}(f_{i-1}, \mathcal{T}^*(h_i))$   
4:  $\mathcal{T}_\theta^\dagger(g) \leftarrow f_l$ 
```


---

# Learned Primal-Dual

$g$


- apply  $\mathcal{T}$
- apply  $\mathcal{T}^*$
- copy
- 3x3 conv + PReLU
- 3x3 conv

# Learned Primal-Dual




- apply  $\mathcal{T}$
- apply  $\mathcal{T}^*$
- copy
- 3x3 conv + PReLU
- 3x3 conv

# Learned Primal-Dual




- apply  $\mathcal{T}$
- apply  $\mathcal{T}^*$
- copy
- 3x3 conv + PReLU
- 3x3 conv

# Learned Primal-Dual




- apply  $\mathcal{T}$
- apply  $\mathcal{T}^*$
- copy
- 3x3 conv + PReLU
- 3x3 conv

# Learned Primal-Dual




- apply  $\mathcal{T}$
- apply  $\mathcal{T}^*$
- copy
- 3x3 conv + PReLU
- 3x3 conv

# Learned Primal-Dual



- apply  $\mathcal{T}$
- apply  $\mathcal{T}^*$
- copy
- 3x3 conv + PReLU
- 3x3 conv


# Learned Primal-Dual




- apply  $\mathcal{T}$
- apply  $\mathcal{T}^*$
- copy
- 3x3 conv + PReLU
- 3x3 conv




# Learned Primal-Dual



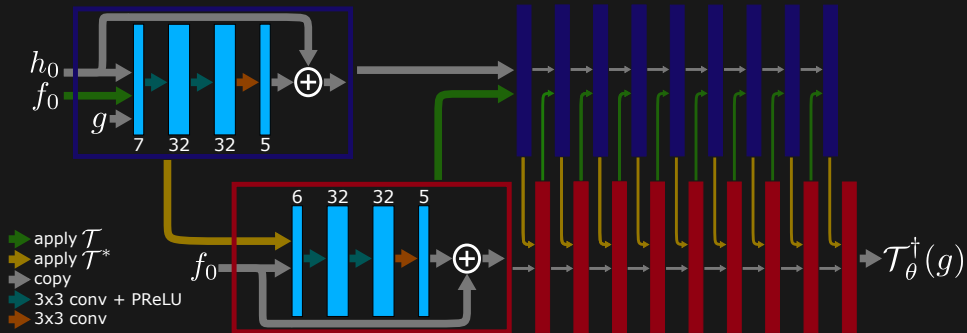
# Learned Primal-Dual








# Learned Primal-Dual








# Learned Primal-Dual








# References

-  *ADMM-Net: A Deep Learning Approach for Compressive Sensing MRI*  
Yang et. al. NIPS 2016
-  *Recurrent inference machines for solving inverse problems*  
Putzky and Welling, arXiv 2017
-  *Learning a Variational Network for Reconstruction of Accelerated MRI Data*  
Hammernick et. al., arXiv 2017
-  *Solving ill-posed inverse problems using iterative deep neural networks*  
Adler and Öktem, Inverse Problems 2017
-  *Learned Primal-Dual Reconstruction*  
Adler and Öktem, IEEE TMI 2018

# References

-  *ADMM-Net: A Deep Learning Approach for Compressive Sensing MRI*  
Yang et. al. NIPS 2016
-  *Recurrent inference machines for solving inverse problems*  
Putzky and Welling, arXiv 2017
-  *Learning a Variational Network for Reconstruction of Accelerated MRI Data*  
Hammernick et. al., arXiv 2017
-  *Solving ill-posed inverse problems using iterative deep neural networks*  
Adler and Öktem, *Inverse Problems* 2017
-  *Learned Primal-Dual Reconstruction*  
Adler and Öktem, IEEE TMI 2018

# References

-  *ADMM-Net: A Deep Learning Approach for Compressive Sensing MRI*  
Yang et. al. NIPS 2016
-  *Recurrent inference machines for solving inverse problems*  
Putzky and Welling, arXiv 2017
-  *Learning a Variational Network for Reconstruction of Accelerated MRI Data*  
Hammernick et. al., arXiv 2017
-  *Solving ill-posed inverse problems using iterative deep neural networks*  
Adler and Öktem, Inverse Problems 2017
-  *Learned Primal-Dual Reconstruction*  
Adler and Öktem, IEEE TMI 2018

# Results

## Results for CT with *Human data*

- ▶ Inverse problem:

$$g = \mathcal{P}(f) + \delta g$$

- ▶ Geometry: fan beam 1000 angles
- ▶ Noise: Poisson noise (low dose CT)
- ▶ Training data: 2000  $512 \times 512$  pixel slices



# Results

## Results for CT with *Human data*


- ▶ Inverse problem:

$$g = \mathcal{P}(f) + \delta g$$


- ▶ Geometry: fan beam 1000 angles
- ▶ Noise: Poisson noise (low dose CT)
- ▶ Training data: 2000  $512 \times 512$  pixel slices

## Compare to:

- ▶ Analytic Pseudo-Inverse (FBP)
- ▶ Variational methods (TV-regularization)
- ▶ Post-processing deep learning by U-Net




Phantom




FBP

PSNR 33.65 dB, SSIM 0.830, 423 ms




Phantom




TV


PSNR 37.48 dB, SSIM 0.946, 64 371 ms




Phantom



FBP + U-Net denoising  
PSNR 41.92 dB, SSIM 0.941, 463 ms



Phantom



Learned Primal-Dual  
PSNR 44.11 dB, SSIM 0.969, 620 ms

# Conclusions

- ▶ Machine learning allows us to handle complicated priors

Source:

[github.com/adler-j](https://github.com/adler-j)

Contact:

[jonasadl@kth.se](mailto:jonasadl@kth.se)

# Conclusions

- ▶ Machine learning allows us to handle complicated priors
- ▶ Fully learned reconstruction is in-feasible

Source:

[github.com/adler-j](https://github.com/adler-j)

Contact:

[jonasadl@kth.se](mailto:jonasadl@kth.se)

# Conclusions

- ▶ Machine learning allows us to handle complicated priors
- ▶ Fully learned reconstruction is in-feasible
- ▶ Learned post-processing gives good results

Source:

[github.com/adler-j](https://github.com/adler-j)

Contact:

[jonasadl@kth.se](mailto:jonasadl@kth.se)



# Conclusions

- ▶ Machine learning allows us to handle complicated priors
- ▶ Fully learned reconstruction is in-feasible
- ▶ Learned post-processing gives good results
- ▶ Learned iterative reconstruction gives better results

Source:

[github.com/adler-j](https://github.com/adler-j)

Contact:

[jonasadl@kth.se](mailto:jonasadl@kth.se)

# Conclusions

- ▶ Machine learning allows us to handle complicated priors
- ▶ Fully learned reconstruction is in-feasible
- ▶ Learned post-processing gives good results
- ▶ Learned iterative reconstruction gives better results

## Questions!

Source:

[github.com/adler-j](https://github.com/adler-j)

Contact:

[jonasadl@kth.se](mailto:jonasadl@kth.se)